

LOCAL AI BENCH 4.1

Local AI Bench Decision Report - commercial-golden-system

COMPLETE EVIDENCE

All 3 recorded performance samples represented here were aggregate-eligible.

Run identity

Field	Value
System	commercial-golden-system
Application	4.1
Engine	llamacpp
OS	Test OS
Backend	Not recorded
Memory	64 GB
Run status	complete
Plan ID	Not recorded
Methodology profile	Legacy / not recorded

Coverage

Stage	Status	Models	Skipped / failed
llm	complete	1	0
conv	complete	1	0
llamabench	complete	1	0

Performance evidence

Workload	Model	Case	Tokens/sec	TTFT sec
Single-shot LLM	golden	2K	50.00	0.250
Conversation	golden	0K	48.00	0.100

Accuracy evidence

Workload	Model	Accuracy	Correct
No evidence recorded			

Sample validity

Workload	Case	Valid	Excluded
Single-shot LLM	golden / 2K	2	0
Conversation	golden / 0K	1	0

Effective optimizations

Recorded setting
No evidence recorded

Interpretation limits

This report presents measured evidence, not a hidden composite score or a universal recommendation. Compare only compatible methodology, model artifacts, runtimes, cache semantics, and effective settings. Missing or invalid data is not zero. Review raw samples, exclusion reasons, and the verified result bundle before making a purchase, launch, or capacity decision.